


MLOps Assignment — Heart Disease Prediction Pipeline	1
1. Overview	2
2. Prerequisites & Local Setup	3
Prerequisites	3
Reproducible setup (Local Setup)	3
Run the API locally.....	4
Deploy the full stack to local Kubernetes.....	4
3. Data Acquisition & EDA	4
Dataset.....	4
Cleaning (dataprocessing/process_data.py).....	4
EDA (notebooks/eda.ipynb)	5
EDA Visualisations	5
4. Feature Engineering & Model Development	7
Preprocessing pipeline (src/data_prep/preprocess.py)	7
Models (src/training/train.py).....	7

Cross-validation results	7
Model selection rationale	7
5. Experiment Tracking.....	8
Experiment tracking summary	8
6. Model Packaging & Reproducibility	9
7. CI/CD Pipeline & Automated Testing	10
Tests (17 total)	10
GitHub Actions (.github/workflows/ci.yml)	10
Manual triggers	11
Local execution with act	11
CI/CD workflow screenshots	11
8. Containerization	12
9. Production Deployment	14
Deployment verification	15
10. Monitoring & Logging	17
Structured request logs	17
Prometheus (pull model)	17
Grafana (auto-provisioned).....	17
11. Conclusion	18

1. Overview

A binary classifier for heart disease, built on the Cleveland UCI dataset, wrapped in a full MLOps pipeline: data processing → training with MLflow tracking → tested FastAPI service → containerized → CI/CD via GitHub Actions → deployed to Docker Desktop Kubernetes with

Prometheus + Grafana monitoring.



2. Prerequisites & Local Setup

Prerequisites

Tool	Version	Purpose
Python	3.12+	Runtime
Docker Desktop	latest	Container runtime + local Kubernetes cluster
kubectl	bundled	Ships with Docker Desktop
Git	any	Clone the repo
act (optional)	latest	Run the GitHub Actions workflow locally (brew install act)

Docker Desktop settings: enable Kubernetes (Settings → Kubernetes → Enable). Allocate at least 4 GB RAM and 2 CPUs (Settings → Resources). MLflow's pod requests up to 1.5 GiB.

macOS users: disable AirPlay Receiver (System Settings → General → AirDrop & Handoff) or accept that MLflow runs on **port 5001** to avoid a conflict on port 5000.

Reproducible setup (Local Setup)

```
git clone https://github.com/Santosh-src/MLOPS_Pipeline.git
cd MLOPS_Pipeline
```

```
python3 -m venv venv && source venv/bin/activate
pip install -r requirements.txt
```

```
python dataprocessing/process_data.py      # cleans the Cleveland raw file
python -m src.training.train                # trains, logs to MLflow, saves mo
```

```
del
pytest tests/ -v # 17 tests
```

Run the API locally

```
uvicorn src.serving.app:app --port 8080
curl http://localhost:8080/health
```

Deploy the full stack to local Kubernetes

```
./scripts/deploy.sh # builds image, applies manifests, waits for rollouts
./scripts/undeploy.sh # tears it all down
```

3. Data Acquisition & EDA

Dataset

Cleveland Heart Disease Database, UCI ML Repository. Original investigators: Robert Detrano, M.D., Ph.D. The full database has 76 attributes; published research uses the standard 14-attribute extraction, which is what we use.

Feature	Description	Type
age	Age in years	Numeric
sex	1 = male, 0 = female	Binary
cp	Chest pain type (1–4)	Categorical
trestbps	Resting blood pressure (mm Hg)	Numeric
chol	Serum cholesterol (mg/dl)	Numeric
fbs	Fasting blood sugar > 120 mg/dl	Binary
restecg	Resting ECG (0–2)	Categorical
thalach	Maximum heart rate achieved	Numeric
exang	Exercise-induced angina	Binary
oldpeak	ST depression induced by exercise	Numeric
slope	Slope of peak exercise ST segment	Categorical
ca	Major vessels coloured by fluoroscopy	Categorical
thal	Thalassemia (3=normal, 6=fixed, 7=rev)	Categorical
target	0 = no disease, 1 = disease (binarised)	Binary

Cleaning (dataprocessing/process_data.py)

The raw UCI file (data/processed.cleveland.data) ships with the repository.

1. Loads with ? mapped to NaN.
2. Imputes numeric columns with median, categorical/binary with mode (6 missing values: 4 in ca, 2 in thal).
3. Binarises target (0 vs 1–4) and casts integer columns.
4. Writes dataprocessing/processed_cleveland.csv (303 rows, 14 cols, 0 missing).


EDA (notebooks/eda.ipynb)

- **Class balance:** 164 no-disease vs 139 disease (54.1% / 45.9%) — workable without resampling.
- **Strongest correlates with target:** thal, ca, oldpeak, exang, cp.
- Disease patients show **lower** max heart rate (thalach) and are slightly older.
- Asymptomatic chest pain (cp = 4) is strongly associated with disease.


EDA Visualisations

Key plots produced in the notebook (also saved to screenshots/):


Class distribution bar chart — 164 no-disease vs 139 disease (54.1%/45.9%), confirming no severe imbalance.





Correlation heatmap — highlights cp, thalach, exang, oldpeak, ca, and thal as the strongest correlates with the target.



Histograms per feature — numeric feature distributions faceted by disease status; disease patients skew older with lower max heart rates.



Box plots — oldpeak and thalach show clear separation between classes.



4. Feature Engineering & Model Development

Preprocessing pipeline (`src/data_prep/preprocess.py`)

A ColumnTransformer packed inside an sklearn Pipeline:

- Numeric (age, trestbps, chol, thalach, oldpeak) → StandardScaler
- Binary (sex, fbs, exang) → passthrough
- Categorical (cp, restecg, slope, ca, thal) → OneHotEncoder(handle_unknown="ignore")

The same pipeline is used at training and inference, so the served model accepts raw feature values.

Models (`src/training/train.py`)

Model	Hyperparameter grid
Logistic Regression	$C \in \{0.01, 0.1, 1.0, 10.0\}$
Random Forest	$n_estimators \in \{50, 100, 200\}$, $max_depth \in \{5, 10, None\}$, $min_samples_split \in \{2, 5\}$

Both are tuned with GridSearchCV(`scoring="roc_auc"`) over a 5-fold StratifiedKFold, then re-evaluated with `cross_validate` across the full metric set.

Cross-validation results

Metric	Logistic Regression	Random Forest
Accuracy	0.8580	0.8415
Precision	0.8874	0.8716
Recall	0.7905	0.7690
F1	0.8343	0.8155
ROC-AUC	0.9192	0.9186

Selected: Logistic Regression ($C = 1.0$). Marginally higher CV ROC-AUC, smaller train/CV gap, faster inference, and easier to explain to a clinician than a forest.

Model selection rationale

Both models achieve strong CV ROC-AUC (>0.91). Logistic Regression was chosen because:

1. **Comparable performance** — only 0.0006 ROC-AUC difference; not statistically significant.
2. **Lower variance** — smaller gap between train and CV metrics, suggesting better generalisation.

3. **Interpretability** — coefficients directly map to feature importance, critical in a clinical setting.
4. **Inference speed** — single matrix multiply vs ensemble of trees; important for low-latency serving.

5. Experiment Tracking

MLflow is wired into `train.py`. Each model gets its own run logging:

- **Params** — model type, best hyperparameters, CV folds
- **Metrics** — accuracy, precision, recall, F1, ROC-AUC for both train set and CV (mean + std)
- **Artifacts** — confusion matrix PNG, ROC curve PNG, the serialized sklearn model

View experiments locally:

```
mlflow ui --port 5001
```

View on Kubernetes (after `./scripts/deploy.sh`):

Open <http://localhost:5001> — the mlops-tracking namespace runs `ghcr.io/mlflow/mlflow:v3.11.1-full` on port 5000 internally, exposed as 5001 to avoid the macOS AirPlay conflict.


Experiment tracking summary

Run	Model	Best Params	CV ROC-AUC	CV Accuracy	CV F1
1	Logistic Regression	C=1.0	0.9192	0.8580	0.8343
2	Random Forest	n_estimators=100, max_depth=None, min_samples_split=2	0.9186	0.8415	0.8155

Artifacts logged per run: confusion matrix (PNG), ROC curve (PNG), serialized sklearn pipeline (joblib), and metrics JSON.

Screenshots:

Experiments:



The screenshot shows the MLflow Experiments web interface. On the left is a sidebar with navigation links: Home, Experiments (selected), and Model registry. The main area displays a table of experiments. At the top right of the main area are buttons for 'Create', 'Compare', and 'Delete'. The table has columns for Name, Time created, Last modified, Description, and Tags. Two experiments are listed: 'heart-disease-classification' and 'Default', both created on 05/10/2026 at 02:50:47 PM and 02:47:57 PM respectively.

Name	Time created	Last modified	Description	Tags
heart-disease-classification	05/10/2026, 02:50:47 PM	05/10/2026, 02:50:47 PM	-	-
Default	05/10/2026, 02:47:57 PM	05/10/2026, 02:47:57 PM	-	-

Runs:

Experiments > heart-disease-classification >

Runs

metrics.rmse < 1 and params.model = "tree"

Time created

State: Active

Datasets

Sort: Created

Columns

+ New run

Expand rows

Group by

	Run Name	Created	Dataset	Duration	User	Source	Version	Models	Description
	RandomForest-grid	1 hour ago	cleveland-train (6e288405) Train... +8	7.6s	sanappan	train.py	aabe8a	model	-
	LogisticRegression...	1 hour ago	cleveland-train (6e288405) Train... +8	15.9s	sanappan	train.py	aabe8a	heart-disease-classifier ...	-

Experiments > heart-disease-classification >

Runs

metrics.rmse < 1 and params.model = "tree"

Time created | State: Active | Datasets | Sort: Created | Columns | + New run

Expand rows | Group by: dataset

Run Name	Created	Metrics								
		accuracy_score	best_cv_score	cv_accuracy	cv_accuracy_s	cv_f1	cv_f1_std	cv_precision	cv_precision_s	cv_recall
Group: dataset.3a04bec2 16	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: cleveland-train.6e288405 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
LogisticRegression-grid	1 hour ago	0.8852459...	0.9038658...	0.8471088...	0.0099769...	0.8245007...	0.0068667...	0.8767529...	0.0541903...	0.783399...
RandomForest-grid	1 hour ago	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: cleveland-test.8134b5b7 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.4622bbce 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.9777fb66 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.0f9cce88 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.bca985f8 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.5c117f86 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...
Group: dataset.8134b5b7 2	-	0.877049...	0.899912...	0.822363...	0.011952...	0.800800...	0.015290...	0.831560...	0.045879...	0.778656...

Registry:

Registered Models >

heart-disease-classifier

Created Time: 05/10/2026, 02:51:03 PM | Last Modified: 05/10/2026, 02:51:03 PM

Description | Edit

Tags

Versions | Compare

New model registry UI

Version	Registered at	Created by	Tags	Aliases	Description
Version 1	05/10/2026, 02:51:03 PM		Add	@ champion	

< Previous | Next > | 25 / page

6. Model Packaging & Reproducibility

- **Artifact:** model/heart_disease_model.joblib — the entire Pipeline (preprocessor + classifier), so inference doesn't need a separate transform step.

- **Metrics summary:** `model/metrics.json` — best model, best params, CV and train metrics.
- **Pinned dependencies:** `requirements.txt`.
- **Determinism:** `random_state = 42` everywhere; cleaning is deterministic given the source file.

```
python3 -m venv venv && source venv/bin/activate
pip install -r requirements.txt
python dataprocessing/process_data.py
python -m src.training.train
```

7. CI/CD Pipeline & Automated Testing

Tests (17 total)

- `tests/test_data.py` — 8 tests on the processed dataset: existence, shape, no nulls, columns, target binary, feature ranges, dtypes, class balance.
- `tests/test_model.py` — 9 tests on the trained model: file exists, loads, has `predict/predict_proba`, prediction shape and binary output, probabilities sum to 1, accuracy ≥ 0.75 , batch prediction.

GitHub Actions ([.github/workflows/ci.yml](https://github.com/workflows/ci.yml))

Three jobs, with sane gating:

ci — runs on every push, PR, or `workflow_dispatch`:

1. Lint (ruff)
2. Process data (download + clean)
3. Data tests (8)
4. Train models (GridSearchCV + MLflow)
5. Model tests (9)
6. Upload artifacts: test results (JUnit XML), trained model, MLflow runs

The data tests must pass *before* training, and the model tests must pass *after* — which catches silent breakage in either layer.

cd — runs after `ci` on push to main (skippable via `skip_cd` input):

1. Pull the trained model artifact
2. Build the Docker image
3. **Container smoke test** — runs the image, hits `/health` and `/predict`. Failing here blocks the push.
4. Tag and push to `ghcr.io` (gated by `push_to_registry` input)
5. If invoked locally via `act`, deploy to Docker Desktop K8s

undeploy — only runs when action=undeploy is selected. Tears the stack down via act (no-op on the GitHub runner, since it can't reach your local cluster).

Manual triggers

Input	Values	Effect
action	deploy / undeploy	Deploy stack (default) or tear it down
skip_cd	false / true	Run only the ci job
push_to_registry	false / true	Push image to ghcr.io (auto-true on push to main)

Local execution with act


```
brew install act
```


```
# CI only – no Kubernetes interaction
act -j ci
```

```
# Anything that touches K8s – wrapper mounts ~/.kube into the act container
./scripts/act-local.sh push
./scripts/act-local.sh workflow_dispatch -j undeploy --input action=undeploy
```

The cd and undeploy jobs install Linux kubectl inside the container and call k8s/act-kubeconfig.sh, which patches the kubeconfig to reach Docker Desktop's K8s API through host.docker.internal with --insecure-skip-tls-verify. The wrapper exists because act does not expand environment variables in .actrc, so the kube mount path can't be made portable there.

CI/CD workflow screenshots





8. Containerization

src/serving/app.py (FastAPI) exposes:

Method	Path	Purpose
GET	/	Liveness
GET	/health	Readiness probe
POST	/predict	Prediction (JSON in/out, Pydantic-validated)
GET	/metrics	Prometheus metrics
GET	/docs	Swagger UI

Dockerfile is a two-stage build on python:3.12-slim: dependencies in stage 1, runtime image in stage 2 with non-root appuser, HEALTHCHECK, and port 8080 exposed.


```
docker build -t heart-disease-api .
docker run -p 8080:8080 heart-disease-api
curl -X POST http://localhost:8080/predict \
  -H "Content-Type: application/json" \
  -d '{"age":63,"sex":1,"cp":3,"trestbps":145,"chol":233,"fbs":1,"restecg":0,
"thalach":150,"exang":0,"oldpeak":2.3,"slope":0,"ca":0,"thal":6}'
```

Sample prediction response:

```
{
  "prediction": 1,
  "confidence": 0.8421,
  "disease_probability": 0.8421
}
```

The /predict endpoint accepts JSON input matching the 13 features, validates via Pydantic, and returns the prediction (0/1), confidence, and disease probability.

Screenshots:



Namespace:

```
sanappan@SANAPPAN-M-FXH7 M-Tech % kubectl get namespace
```

NAME	STATUS	AGE
default	Active	11d
kube-node-lease	Active	11d
kube-public	Active	11d
kube-system	Active	11d
mlops-app	Active	70m
mlops-monitoring	Active	70m
mlops-tracking	Active	70m

Pods:

```
sanappan@SANAPPAN-M-FXH7 M-Tech % kubectl get all -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	pod/coredns-668d6bf9bc-5lwsn	1/1	Running	3 (2d7h ago)	11d
kube-system	pod/coredns-668d6bf9bc-s2wrt	1/1	Running	3 (2d7h ago)	11d
kube-system	pod/etcd-docker-desktop	1/1	Running	4 (2d7h ago)	11d
kube-system	pod/kube-apiserver-docker-desktop	1/1	Running	4 (2d7h ago)	11d
kube-system	pod/kube-controller-manager-docker-desktop	1/1	Running	4 (2d7h ago)	11d
kube-system	pod/kube-proxy-xnsmb	1/1	Running	4 (2d7h ago)	11d
kube-system	pod/kube-scheduler-docker-desktop	1/1	Running	3 (2d7h ago)	11d
kube-system	pod/storage-provisioner	1/1	Running	9 (3h43m ago)	11d
kube-system	pod/vpnkit-controller	1/1	Running	5 (2d7h ago)	11d
mlops-app	pod/heart-disease-api-55b967cc7c-q7nzm	1/1	Running	0	70m
mlops-app	pod/heart-disease-api-55b967cc7c-zlh9m	1/1	Running	0	70m
mlops-monitoring	pod/grafana-5d766986f-9nqtc	1/1	Running	1 (65m ago)	70m
mlops-monitoring	pod/prometheus-5bcf687b65-dcdrc	1/1	Running	0	70m
mlops-tracking	pod/mlflow-server-5cd6c79447-j5tkq	1/1	Running	0	70m

Service:

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	11d
kube-system	service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP, 9153/TCP	11d
mlops-app	service/heart-disease-api	LoadBalancer	10.102.59.228	localhost	8080:31955/TCP	70m
mlops-monitoring	service/grafana	LoadBalancer	10.109.183.141	localhost	3000:32632/TCP	70m
mlops-monitoring	service/prometheus	LoadBalancer	10.96.23.93	localhost	9090:31603/TCP	70m
mlops-tracking	service/mlflow-server	LoadBalancer	10.99.243.148	localhost	5001:31816/TCP	70m

Deployment & Replicas:

NAMESPACE	NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-system	daemonset.apps/kube-proxy	1	1	1	1	1	kubernetes.io/os=linux	11d
NAMESPACE	NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	AGE	
kube-system	deployment.apps/coredns	2/2	2	2	2	2	11d	
mlops-app	deployment.apps/heart-disease-api	2/2	2	2	2	2	70m	
mlops-monitoring	deployment.apps/grafana	1/1	1	1	1	1	70m	
mlops-monitoring	deployment.apps/prometheus	1/1	1	1	1	1	70m	
mlops-tracking	deployment.apps/mlflow-server	1/1	1	1	1	1	70m	
NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE			
kube-system	replicaset.apps/coredns-668d6bf9bc	2	2	2	11d			
mlops-app	replicaset.apps/heart-disease-api-55b967cc7c	2	2	2	70m			
mlops-monitoring	replicaset.apps/grafana-5d766986f	1	1	1	70m			
mlops-monitoring	replicaset.apps/prometheus-5bcf687b65	1	1	1	70m			
mlops-tracking	replicaset.apps/mlflow-server-5cd6c79447	1	1	1	70m			

9. Production Deployment

Deployed to Docker Desktop's built-in Kubernetes across three namespaces.

Namespace	Components	Notes
mlops-app	Deployment (2 replicas) + Service + Ingress	Probes, resource limits, imagePullPolicy: Never
mlops-tracking	MLflow Deployment + Service	Official ghcr.io/mlflow/mlflow:v3.11.1-full image, SQLite backend, single worker, 1.5 GiB limit
mlops-monitoring	Prometheus + Grafana Deployments + Services	Cross-namespaces scraping; Grafana auto-provisioned via ConfigMaps

Manifests (k8s/): namespaces.yaml, deployment.yaml, service.yaml, ingress.yaml, mlflow.yaml, prometheus.yaml, grafana.yaml.

One-shot deploy:

```
./scripts/deploy.sh
```

Endpoint	URL
API	http://localhost:8080
API (Ingress, optional)	http://heart-disease.local
MLflow UI	http://localhost:5001
Prometheus	http://localhost:9090
Grafana	http://localhost:3000 (admin / admin)

Tear down:

```
./scripts/undeploy.sh           # delete the three namespaces
./scripts/undeploy.sh --all      # also remove the local heart-disease-api D
ocker image
```

Deployment verification


After `./scripts/deploy.sh` completes, all pods are running:

NAMESPACE	NAME	READY	STATUS	RE
STARTS				
mlops-app	heart-disease-api-xxx-yyy	1/1	Running	0
mlops-app	heart-disease-api-xxx-zzz	1/1	Running	0
mlops-tracking	mlflow-xxx-yyy	1/1	Running	0
mlops-monitoring	prometheus-xxx-yyy	1/1	Running	0
mlops-monitoring	grafana-xxx-yyy	1/1	Running	0


Endpoints are verified with curl:

```
curl http://localhost:8080/health # {"status": "healthy", "model_Loaded": true}
curl http://localhost:5001 # MLflow UI accessible
curl http://localhost:9090/targets # Prometheus shows heart-disease-api target UP
curl http://localhost:3000 # Grafana dashboard loads
```


Screenshots:




/metrics




/health



Predict 1:



Predict 2:



10. Monitoring & Logging

Structured request logs

Every prediction emits a single JSON line on stdout:

```
{"event": "prediction", "input": {...}, "prediction": 1, "disease_probability": 0.84, "confidence": 0.84, "latency_ms": 3.21}
```

Picked up by `kubectl logs` and any log shipper.

Prometheus (pull model)

`prometheus-fastapi-instrumentator` adds `/metrics` to the FastAPI app. Prometheus scrapes it every 15 seconds at:

```
heart-disease-api.mlops-app.svc.cluster.local:8080/metrics
```

Standard metrics exposed: `http_requests_total`, `http_request_duration_seconds`, `http_requests_in_progress`.

Grafana (auto-provisioned)


Defined in `k8s/grafana.yaml` via three ConfigMaps:

- **Datasource** — Prometheus at the in-cluster service URL, marked default.
- **Dashboard provider** — file provisioner pointing at `/var/lib/grafana/dashboards`.
- **Dashboard** — “Heart Disease API Monitoring” with six panels: request rate, latency p50/p95/p99, requests by status, error rate, health status, in-flight requests.


After `./scripts/deploy.sh`, open <http://localhost:3000> (admin/admin) — the dashboard is already there.

Screenshots:


Prometheus:



Grafana:



Metrics:



11. Conclusion

The pipeline covers the full lifecycle end to end:

- Deterministic data processing with download + cleaning
 - Two models compared with cross-validation; Logistic Regression selected
 - MLflow tracking of params, metrics, and plots for every run
 - 17 unit tests gating CI; container smoke test gating CD
 - Multi-stage Docker image, non-root, healthchecked
 - Kubernetes deployment across three namespaces with auto-provisioned monitoring
 - The whole CI/CD workflow is runnable locally with act (including real K8s deploy and undeploy)
-